

Bilevel optimization: A functional perspective

Application to self-supervised learning

MS386

SIAM OP26

Michael Arbel

June 05, 2026

Functional Bilevel Optimization

Definition

Let $\mathcal{U}$ and $\mathcal{H}$ be Hilbert spaces and $L, G : \mathcal{U} \times \mathcal{H} \rightarrow \mathbb{R}$ two smooth functionals, with $P \mapsto G(U, P)$ strongly convex for all $U \in \mathcal{U}$. Consider the problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

s.t.

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Functional Bilevel Optimization

Definition

Let $\mathcal{U}$ and $\mathcal{H}$ be Hilbert spaces and $L, G : \mathcal{U} \times \mathcal{H} \rightarrow \mathbb{R}$ two smooth functionals, with $P \mapsto G(U, P)$ strongly convex for all $U \in \mathcal{U}$. Consider the problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

s.t.

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Why "functional" ?

1. Some problems are naturally defined in function spaces
(i.e. $\mathcal{U}$ and/or $\mathcal{H}$ infinite-dimensional spaces of functions)



Function space analysis reveals the target objects (properties, stability, etc)

2. Reducing the problem to finite dimensional parameters can, sometimes, be misleading

Functional Bilevel Optimization

Definition

Let $\mathcal{U}$ and $\mathcal{H}$ be Hilbert spaces and $L, G : \mathcal{U} \times \mathcal{H} \rightarrow \mathbb{R}$ two smooth functionals, with $P \mapsto G(U, P)$ strongly convex for all $U \in \mathcal{U}$. Consider the problem


$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

s.t.

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Causal effect estimation


Goal: Estimate treatment effect



Causal model **constrains** prediction

Data-driven PDE discovery


Goal: Learn unknown PDE



Variational PDE **constrains** the state

Bilevel Self-supervised learning

Goal: Learn data encoder from views



Regression **constrains** prediction

Functional Bilevel Optimization

Definition

Let $\mathcal{U}$ and $\mathcal{H}$ be Hilbert spaces and $L, G : \mathcal{U} \times \mathcal{H} \rightarrow \mathbb{R}$ two smooth functionals, with $P \mapsto G(U, P)$ strongly convex for all $U \in \mathcal{U}$. Consider the problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

s.t.

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

- L and G might depend on data (learning problem).
- $\mathcal{U}$ and $\mathcal{H}$ might be infinite dimensional spaces of functions
- Elements $U \in \mathcal{U}$ and $P \in \mathcal{H}$ might be functions of high-dimensional vectors

How to solve FBO in practice?

Need function approximation
e.g., neural networks

I- Functional Implicit Differentiation

Leveraging neural networks and gradient methods for solving FBO


FBO problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U) \quad \text{s.t.} \quad P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

P is a function of high-dimensional input x
 $\mathcal{H}$ is infinite dimensional space of functions

How?

Neural approximation



$$f_{\psi}(\cdot) \approx P_U(\cdot)$$


Leveraging neural networks and gradient methods for solving FBO

FBO problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U) \text{ s.t. } P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

P is a function of high-dimensional input x
 $\mathcal{H}$ is infinite dimensional space of functions

How?



Neural approximation

$$f_\psi(\cdot) \approx P_U(\cdot)$$

Differentiable programming

Solve surrogate problem
with wrong implicit object $\psi^\star(\theta)$

$$\begin{aligned} & \min_{U \in \mathcal{U}} L(U, f_{\psi^\star(U)}) \\ \text{s.t. } & \psi^\star(U) = \arg \min_{\psi \in \Psi} G(U, f_\psi) \end{aligned}$$



Multiple choices for the solution $\psi^\star(U)$



Ill-defined surrogate problem

Penalty methods

Blur objective and constraint

$$\min_{U \in \mathcal{U}, \psi \in \Psi} L(U, f_\psi) + \lambda \|\nabla G(U, f_\psi)\|$$

f_ψ can fit both L and C for arbitrary U



U is no longer learned

Neural operators

Replace the full solution map

$$U \mapsto P_U$$

by neural approximation

Tries to solve a harder problem



Data inefficient and costly

The curse of network overparametrization

The curse of dimensions

FuncBO (NeurIPS 2024)

Functional Bilevel Optimization for Machine Learning

Ieva Petrulionyte, Julien Mairal, Michael Arbel

Univ. Grenoble Alpes, Inria, CNRS, Grenoble INP, LJK, 38000 Grenoble, France

`firstname.lastname@inria.fr`



Follow up: Petrulionyte et al, [Non-Stationary Functional Bilevel Optimization](#), AISTATS 2026

FuncBO: Differentiate first, then approximate

FBO problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

s.t.

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Standard differentiable programming approach

Neural approximation *of inner solution P*

yields surrogate problem

Differentiation
in parameter space *of surrogate objective*



Ill-defined surrogate problem

FunBO's approach: Functional Implicit Differentiation

Differentiation
in functional space *of FBO objective*

yields functional gradient

Neural approximation *of gradient components
(e.g., predictor sensitivity)*



Well-defined ICL gradient

Implicit differentiation in L_2 spaces and adjoint sensitivity method

- $\mathcal{U}$ is a Euclidean space, $\mathcal{H}$ is space of square integrable functions under data distribution $\mathbb{P}$
- Inner and outer objectives are expectations of point-wise losses under a data distribution $\mathbb{P}$

$$L(U, P) = \mathbb{E}[\ell(U, P(x), y)]$$

$$G(U, P) = \mathbb{E}[g(U, P(x), y)]$$

Proposition (informal)

If $p \mapsto g(U, p, y)$ is strongly convex + regularity, integrability and growth assumptions:

$$\nabla J(U) = F(U, P_U, a_U) \triangleq \mathbb{E} \left[\partial_1 \ell(U, P_U(x), y) + \partial_{1,2} g(U, P_U, (x), y) a_U(x) \right]$$

where the a_U is the unique minimizer in $L_2(\mathbb{P})$ of the objective: $a \mapsto L_{adj}(U, P_U, a)$

$$L_{adj}(U, P, a) = \mathbb{E} \left[\frac{1}{2} a(x)^\top \partial_{2,2} g(U, P(x), y) a(x) + \partial_2 \ell(U, P(x), y) a(x) \right]$$

The FuncBO algorithm

FBO problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Differentiation in functional space

$$\nabla J(U) = F(U, P_U, a_U)$$

$$a_U = \arg \min_{a \in \mathcal{H}} L_{adj}(U, P_U, a)$$

Neural approximation

$$f_\psi \approx \arg \min_{f_\psi \in NN} \hat{G}(U, f_\psi) \approx P_U$$

$$g_\theta \approx \arg \min_{g_\theta \in NN} \hat{L}_{adj}(U, f_\psi, g_\theta) \approx a_U$$

$$\hat{F}(U, f_\psi, g_\theta) \approx \nabla J(U)$$

- $\hat{F}$, $\hat{G}$ and $\hat{L}_{adj}$ are empirical averages to estimate F , G and L_{adj}
- f_ψ and g_θ obtained by gradient descent w.r.t the parameters ψ and θ
- U updated iteratively using gradient estimate $\hat{F}(U, f_\psi, g_\theta)$

The FuncBO algorithm

FBO problem

$$\min_{U \in \mathcal{U}} J(U) = L(U, P_U)$$

$$P_U = \arg \min_{P \in \mathcal{H}} G(U, P)$$

Differentiation in functional space

$$\nabla J(U) = F(U, P_U, a_U)$$

$$a_U = \arg \min_{a \in \mathcal{H}} L_{adj}(U, P_U, a)$$

Neural approximation

$$f_\psi \approx \arg \min_{f_\psi \in NN} \hat{G}(U, f_\psi) \approx P_U$$


$$g_\theta \approx \arg \min_{g_\theta \in NN} \hat{L}_{adj}(U, f_\psi, g_\theta) \approx a_U$$

$$\hat{F}(U, f_\psi, g_\theta) \approx \nabla J(U)$$

- **Memory and time savings:** No need for higher-order derivatives w.r.t network parameters
- **Stability:** Strongly-convex adjoint objective in function space $\longrightarrow$ Well-defined gradient

Application to causal effect estimation


Results on synthetic task from [Xu et al. 2020]



- Large improvement over classical BO algorithms (AID/ITD)
- Competitive with problem specific method (DFIV)

Causal effect estimation

Fit U using variation induced by Z



$$\min_{U \in \mathcal{U}} \mathbb{E} [\|Y - P_U(Z)\|^2] \leftarrow L(U, P_U)$$

Instrument-induced
Projection P_U

Project $g_U(A)$ onto functions of Z

$$P_U = \arg \min_{P \in L_2(\mathbb{P})} \mathbb{E} [\|P(Z) - g_U(A)\|^2]$$

$G(U, P)$

Learning theory for FBO

- Can we get convergence/generalization results for FuncBO?
- NeurIPS 2025: first attempt when using kernel methods for function approximation

Learning Theory for Kernel Bilevel Optimization

Fares El Khoury¹ Edouard Pauwels² Samuel Vaiter³ Michael Arbel¹




- Next step: dealing with neural network approximation (coming soon)

II- Self-supervised learning as functional bilevel optimization


Self-supervised learning

Learn data encoder without supervision by comparing altered/partial views of the same input

Non-Contrastive SSL



Downstream task




Self-distillation methods for SSL: A puzzle

- State-of-the-art performance
- Can collapse to trivial representation
i.e: $U = V = cst, P = Id$
- Uses heuristics to avoid collapse (Stop-gradient, EMA)
- Not an optimization algorithm [Ponce et al, 2026]

Goal: Understanding:

- What is learned?
- When does collapse occur?
- How to reliably prevent collapse? Bilevel optimization to the rescue!



Self-distillation SSL: Existing theory is informative but restricted

Special setting

- Linear encoders U, V
- Particular data covariance structure
- Special initialization (ex: orthogonal)

Tian et al. 2021

Littwin et al. 2024

Liu et al. 2022

Tang et al. 2023

?

General setting

- Non-linear encoders U, V
- General data distributions
- General initialization

PEIRA: Learning Predictive Encoders through Inter-View Regressor Alignment

Michael Arbel*


Univ. Grenoble Alpes, Inria
CNRS, Grenoble INP, LJK
michael.arbel@inria.fr

Basile Terver*

Ecole Normale Supérieure / PSL
Inria Paris
basile.terver@ens.psl.edu

Jean Ponce

Ecole Normale Supérieure / PSL
New York University
jean.ponce@ens.fr



- Characterization of a self-distillation dynamics for SSL
 - Limiting objects, convergence of the dynamics and characterization of collapse
- Bilevel optimization alternative to SSL that provably avoids collapse

A self-distillation dynamics in function space

$$\mathcal{L}(P, U, V) = \frac{1}{2} \mathbb{E}[\|PU(x) - V(y)\|^2] + \frac{\lambda}{2} \|P\|^2 + \frac{\lambda}{2} \mathbb{E}[\|U(x)\|^2]$$

$$\begin{cases} P^\star = \arg \min_P \mathcal{L}(P, U, V) \\ \dot{U} = -\partial_U \mathcal{L}(P^\star, U, V) \\ V = U \end{cases}$$

- Optimal linear predictor (as in Tang et al. 2023)
- No parametric form for the encoder U
- Dynamics defined in function space $L_2(\mathbb{P}_x, \mathbb{R}^k)$

A self-distillation dynamics in function space: Convergence

$$\mathcal{L}(P, U, V) = \frac{1}{2}\mathbb{E}[\|PU(x) - V(y)\|^2] + \frac{\lambda}{2}\|P\|^2 + \frac{\lambda}{2}\mathbb{E}[\|U(x)\|^2]$$

$$\begin{cases} P^\star = \arg \min_P \mathcal{L}(P, U, V) \\ \dot{U} = -\partial_U \mathcal{L}(P^\star, U, V) \\ V = U \end{cases}$$

- Optimal linear predictor (as in Tang et al. 2023)
- No parametric form for the encoder U
- Dynamics defined in function space $L_2(\mathbb{P}_x, \mathbb{R}^k)$

Theorem (Informal): Assume x and y play symmetric roles and data regular enough.


Then: 1- For any initialization U_0 , the dynamics converges to an equilibrium $U^\star$.

2- The following objective decreases in time (Lyapunov function):

$$\mathcal{F}(U) = -\frac{1}{2}\text{Tr}(\mathbb{E}[U(x)U(y)^\top]^2) + \frac{1}{3}\text{Tr}((\mathbb{E}[U(x)U(x)^\top] + \lambda I)^3)$$

No assumption on the data covariances, or orthogonal initialization, as in Tang et al. 2023

A self-distillation dynamics in function space: Convergence



- No limit cycles, no divergence.
- Convergence to a critical point of the Lyapunov function.

Key technical ingredients

- **Gradient-like system:** The vector field aligns well with the gradient of the Lyapunov function
- **Lojasiewicz-Simon inequality:** Small gradients $\implies$ Small gap to a critical value.
- **Trapping argument:** The trajectory is trapped near an accumulation point (Bolte et al. 2014)
- **Existence of accumulation points:** Non-trivial in infinite dimensions (Palais & Smale 1964)

A self-distillation dynamics in function space: Equilibria

What are the equilibria? Which ones are stable?

Canonical Correlation Analysis to the rescue

Non-Linear Canonical Correlation Analysis (Michaeli et al. 2015)

Correlation maximization

$$\max_{(U,V) \in \mathcal{U} \times \mathcal{V}} \mathbb{E}[U(x)^\top V(y)]$$

$L_2(\mathbb{P}_x)^k$

$L_2(\mathbb{P}_y)^k$

s.t.

$$\mathbb{E}[U(x)U(x)^\top] = I$$

$$\mathbb{E}[V(y)V(y)^\top] = I$$

whitening constraints

Non-Linear Canonical Correlation Analysis (Michaeli et al. 2015)

Correlation maximization

$$\max_{(U,V) \in \mathcal{U} \times \mathcal{V}} \mathbb{E}[U(x)^\top V(y)]$$

$L_2(\mathbb{P}_x)^k$

$L_2(\mathbb{P}_y)^k$

s.t.

$$\mathbb{E}[U(x)U(x)^\top] = I$$

$$\mathbb{E}[V(y)V(y)^\top] = I$$

whitening constraints

Non-linear CCA solution: top-k canonical functions

$$U(x) = \begin{pmatrix} u_1(x) \\ \vdots \\ u_k(x) \end{pmatrix}, \quad V(y) = \begin{pmatrix} v_1(y) \\ \vdots \\ v_k(y) \end{pmatrix}$$

Canonical correlation $1 = s_1 \geq s_2 \geq \dots$

$$\mathbb{E}[u_i(x)v_i(y)] = s_i \qquad \mathbb{E}[u_i(x)v_j(y)] = 0$$

De-correlation (when $i \neq j$)

Orthonormal basis of $L_2(\mathbb{P}_y)$

Orthonormal basis of $L_2(\mathbb{P}_x)$

SimSiam-like dynamics in function space: Equilibria and Stability

Theorem (Informal):

Equilibria span CCA subspaces. The only stable ones are of the form:

$$U^\star(x) = \sum_{i=1}^r f(s_i, \lambda) q_i u_i(x)$$

Diagram illustrating the components of the equilibrium form:

- Any $0 \leq r \leq k$** : Points to the index r in the summation.
- Regularization**: Points to the parameter λ in the function $f(s_i, \lambda)$.
- Any orthonormal family $(q_i)_{1 \leq i \leq r}$ of $\mathbb{R}^k$** : Points to the vector q_i .
- Canonical function**: Points to the function $u_i(x)$.
- Canonical correlation $s_1 \geq s_2 \geq \dots$** : Points to the scalar s_i .
- Spectral filter**: Points to the function $f(s, \lambda)$.

The spectral filter function is defined as:

$$f(s, \lambda) = \frac{1}{\sqrt{2}}(s + \sqrt{s^2 - 4\lambda}) \mathbf{1}_{s^2 \geq 4\lambda}$$

Collapse occurs when an admissible CCA direction ($f(s_i, \lambda) > 0$ for $i \leq k$) is not selected ($r < i$)

**Preventing collapse through
bilevel optimization**

Intuition and Motivation: Trace maximization of a predictor

- Optimal predictor between two views:

The diagram shows the formula for the optimal predictor $P^\star$ with three annotations:

- Shared signal:** A red arrow points to the matrix $\mathbb{E} [V(y)U(x)^\top]$, which is highlighted in a purple box.
- Within view variability (noise):** A blue arrow points to the matrix $\mathbb{E} [U(x)U(x)^\top]$, which is highlighted in a blue box.
- Signal to noise ratio (matrix form):** A green arrow points to the entire term $\left(\mathbb{E} [U(x)U(x)^\top] + \lambda I \right)^{-1}$, which is highlighted in a green box.

$$P^\star = \mathbb{E} [V(y)U(x)^\top] \left(\mathbb{E} [U(x)U(x)^\top] + \lambda I \right)^{-1}$$

Intuition and Motivation: Trace maximization of a predictor

- Optimal predictor between two views:

The diagram shows the formula for the optimal predictor $P^\star$ with three annotations:

- Shared signal:** A red arrow points to the matrix $\mathbb{E} [V(y)U(x)^\top]$, which is highlighted in a purple box.
- Within view variability (noise):** A blue arrow points to the matrix $\mathbb{E} [U(x)U(x)^\top]$, which is highlighted in a blue box.
- Signal to noise ratio (matrix form):** A green arrow points to the entire term $\left(\mathbb{E} [U(x)U(x)^\top] + \lambda I \right)^{-1}$, which is highlighted in a green box.

$$P^\star = \mathbb{E} [V(y)U(x)^\top] \left(\mathbb{E} [U(x)U(x)^\top] + \lambda I \right)^{-1}$$

$Tr(P^\star)$ acts as a scalar summary of the matrix signal to noise ratio

Predictive Embeddings through Inter-View Regressor Alignment

- Find optimal symmetrized regressor for each encoders pairs (U, V)

$$\frac{1}{2}\mathbb{E}[\|PU(x) - V(y)\|^2] + \frac{1}{2}\mathbb{E}[\|PV(y) - U(x)\|^2] + \frac{\lambda}{2}\|P\|^2$$

$$P_{U,V}^\star = \arg \min_{P \in \mathbb{R}^{k \times k}} \mathcal{L}_{reg}(P; U, V)$$

- Learn encoders by minimizing regularized negative trace of optimal regressor:

Negative trace minimization to align representation

Regularization to fix the scale

$$\mathcal{L}_{PEIRA}(U, V) = -\frac{1}{2}\text{Tr}(P_{U,V}^\star) + \frac{\lambda}{2}\mathbb{E}[\|V(x)\|^2 + \|U(y)\|^2]$$

PEIRA objective: Global maximizers

Theorem (Informal): Assume the data distribution is regular enough (existence of non-linear CCA).

Then: 1- Global minimizers of $\mathcal{L}_{PEIRA}$ are exactly:

$$(U^\star(x), V^\star(y)) = \sum_{i=1}^k f(s_i, \lambda) q_i (u_i(x), v_i(y))$$

Encoder's dimension $\rightarrow k$


Any orthonormal family $(q_i)_{1 \leq i \leq r}$ of $\mathbb{R}^k$

Canonical functions $(u_i(x), v_i(y))$

Soft-thresholding filter $\frac{1}{\sqrt{2}} \left(\max(\sqrt{s} - \lambda), 0 \right)^{\frac{1}{2}}$

$f(s_i, \lambda)$


PEIRA objective: Stability and non-collapse




Theorem (Informal): Under the same assumptions as before:

- 1- Only global minimizers are stable.
- 2- All other critical points are unstable.

Optimizing PEIRA objective: A "bad" Algorithm



Optimizing PEIRA objective: A better Algorithm



Proposition:

$\mathcal{L}_{PEIRA}(U, V)$ has the same gradient as that of an auxiliary loss $\mathcal{L}_{aux}(U, V; P, Q)$ if:

- P is the optimal regressor
- Q satisfies

$$Q = (\mathbb{E} [U(x)U(x)^\top + V(y)V(y)^\top] + \lambda I)^{-1}$$

The auxiliary loss:


$$\begin{aligned} \mathcal{L}_{aux}(U, V; P, Q) = & \frac{1}{2} \mathbb{E} \left[QU(x)^\top (PU(x) - V(y)) \right] \\ & + \frac{1}{2} \mathbb{E} \left[QV(y)^\top (PV(y) - U(x)) \right] \\ & + \frac{\lambda}{2} \mathbb{E} [\|U(x)\|^2 + \|V(y)\|^2] \end{aligned}$$

Pretraining on ImageNet1k and Cifar10

Method	Approach	ImageNet-1K $\uparrow$	CIFAR-10 $\uparrow$
BYOL [Grill et al., 2020]	self-distillation (EMA)	69.3	–
SimSiam [Chen and He, 2021]	self-distillation (stop-gradient)	67.9	–
SimCLR [Chen et al., 2020]	contrastive (negatives)	68.5	–
VICReg [Bardes et al., 2022]	optimization	68.81 ± 0.09	90.92 ± 0.14
SIGReg [Balestriero and LeCun, 2025]	optimization	66.26 ± 0.28	92.34 ± 0.21
PEIRA (ours)	optimization	66.50 ± 0.02	90.97 ± 0.10

Competitive with some SOTA methods

SSL for images: Cifar10



Strong correlation between Val acc and PEIRA objective

Conclusion: Part II

Summary

- There is a formal connexion between self-distillation SSL methods and non-linear CCA
- Collapse is still an issue for these methods
- PEIRA leverages this connection with CCA while provably avoiding collapse
- Promising results on image representation

Future directions/open questions

- More applications? Video, time series, world-models
- Smarter algorithms? Avoiding full matrix inversion? Variance reduction, etc
- Statistical guarantees?

Thank you!